


Hash-Based Indexing



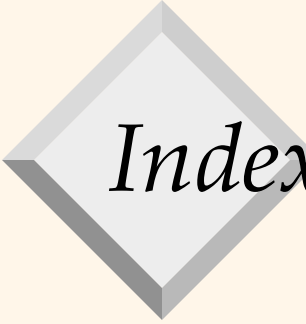
Tree Indexing Summary

- ❑ Static and dynamic data structures
 - ISAM and B+ trees
- ❑ Speed up both range and equality searches
- ❑ B+ trees very widely used in practice
- ❑ ISAM trees can be useful if dataset relatively static (not a lot of overflow pages)
 - Because index is static, no need to lock index pages



Indexing using Hashing

- ❑ Hash-based indexes are for *equality selections*. **Cannot** support range searches.
- ❑ Static and dynamic hashing techniques exist; trade-offs similar to ISAM vs. B+ trees.



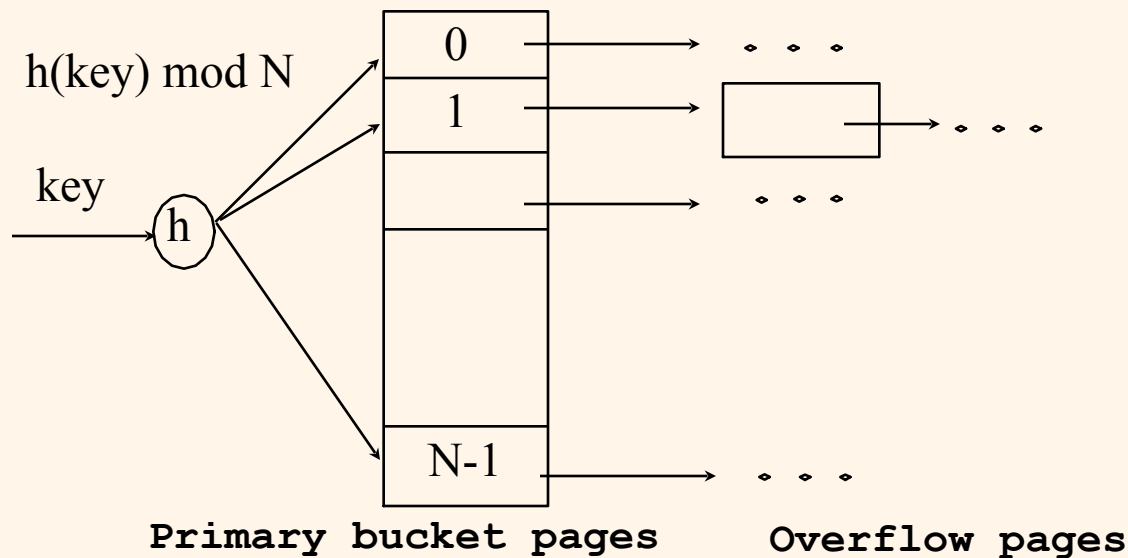
Indexing Using Hashing

[?] *As for any index, 3 alternatives for data entries \mathbf{k}^* :*

- Data record with key value \mathbf{k}*
- $\langle \mathbf{k}, \text{rid of data record with search key value } \mathbf{k} \rangle$*
- $\langle \mathbf{k}, \text{list of rids of data records with search key } \mathbf{k} \rangle$*

Static Hashing


- ❑ For each key k , compute some *hash function* $h(k)$
- ❑ $h(k) \bmod N =$ bucket to which data entry with key k belongs. ($N = \#$ of buckets)





Static Hashing

- ☐ Hash fn works on *search key* field of record r .
- $h(k) \bmod N$ must distribute values over range $0 \dots N-1$.
 - $h(k) = (a * k + b)$ usually works well.
 - a and b are constants that can be used to tune h



Static Hashing

- ☐ Primary bucket pages fixed, allocated sequentially, never de-allocated; overflow pages if needed
 - If no overflow pages, lookup just one disk I/O
 - Overflow pages degrade performance



Static Hashing

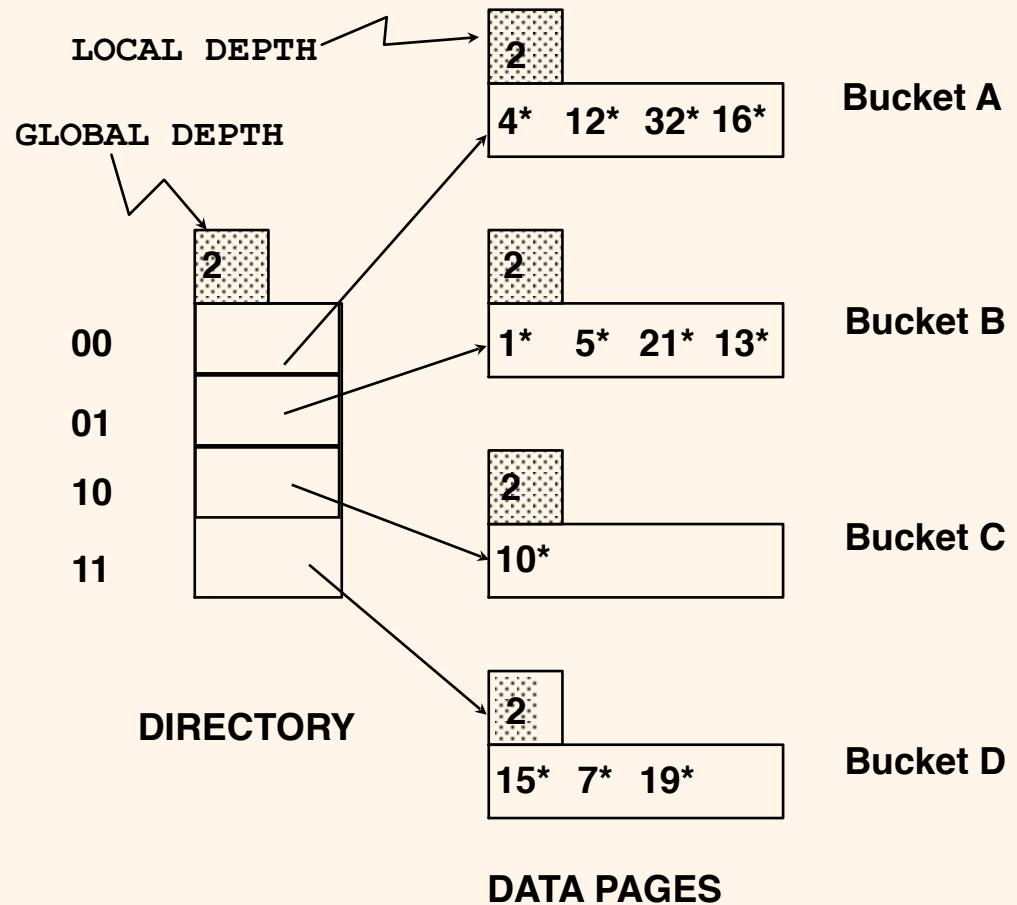
- ❑ If too many overflow pages, can **rehash** (reorganize into more buckets)
 - E.g. double the number of buckets
- ❑ Problems:
 - Takes time if index is big
 - Index cannot be used during rehashing



Extendible Hashing

- ❑ Better idea: split only the bucket that has overflowed
- ❑ Keep a *directory* of pointers to buckets, so searches for the old bucket are properly directed to the two new ones
 - When bucket is split, only directory needs to be adjusted, and this is pretty small

Example



- ❑ Notation: 1* = data entry with *hash value 1*
- ❑ To find bucket for record with key k , look at last 2 bits of $h(k)$
- ❑ E.g. if $h(k) = 14$ i.e. 1110, goes into 3rd bucket

Example

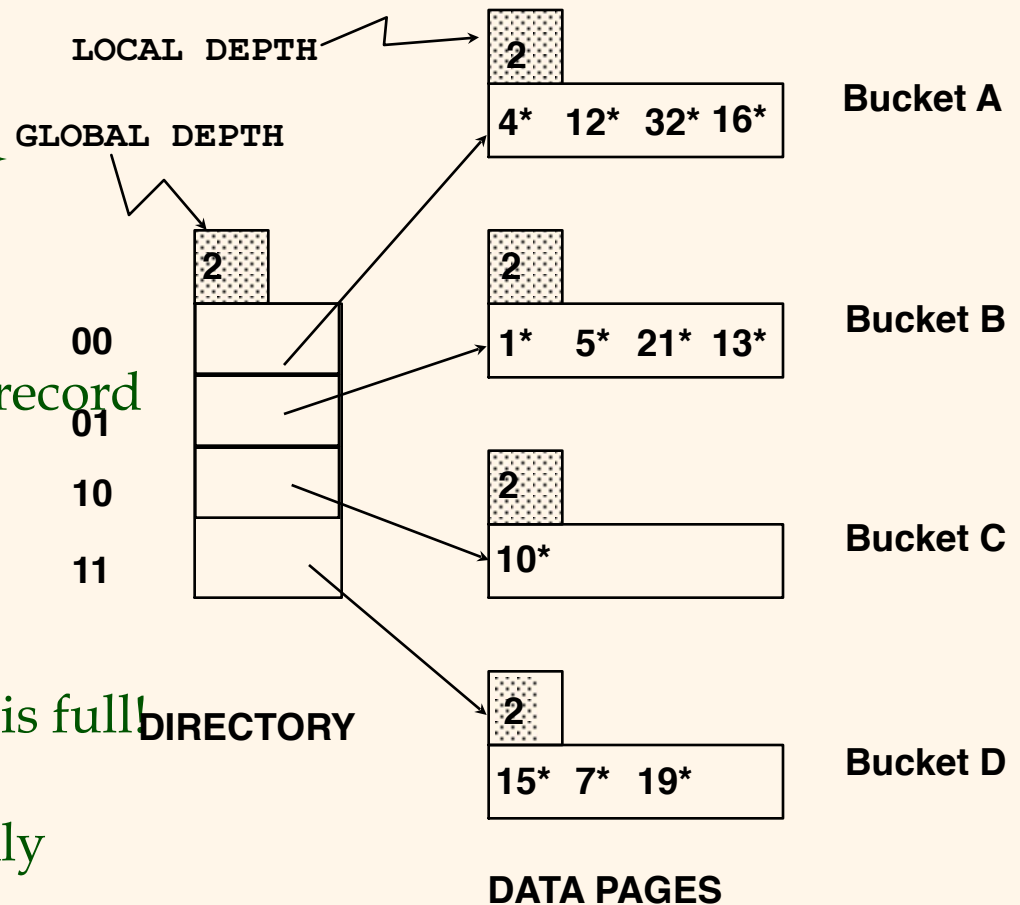
❖ Now suppose want to insert record with $h(k) = 20$

❖ 20 in binary is 10100

❖ Belongs in first bucket which is full!

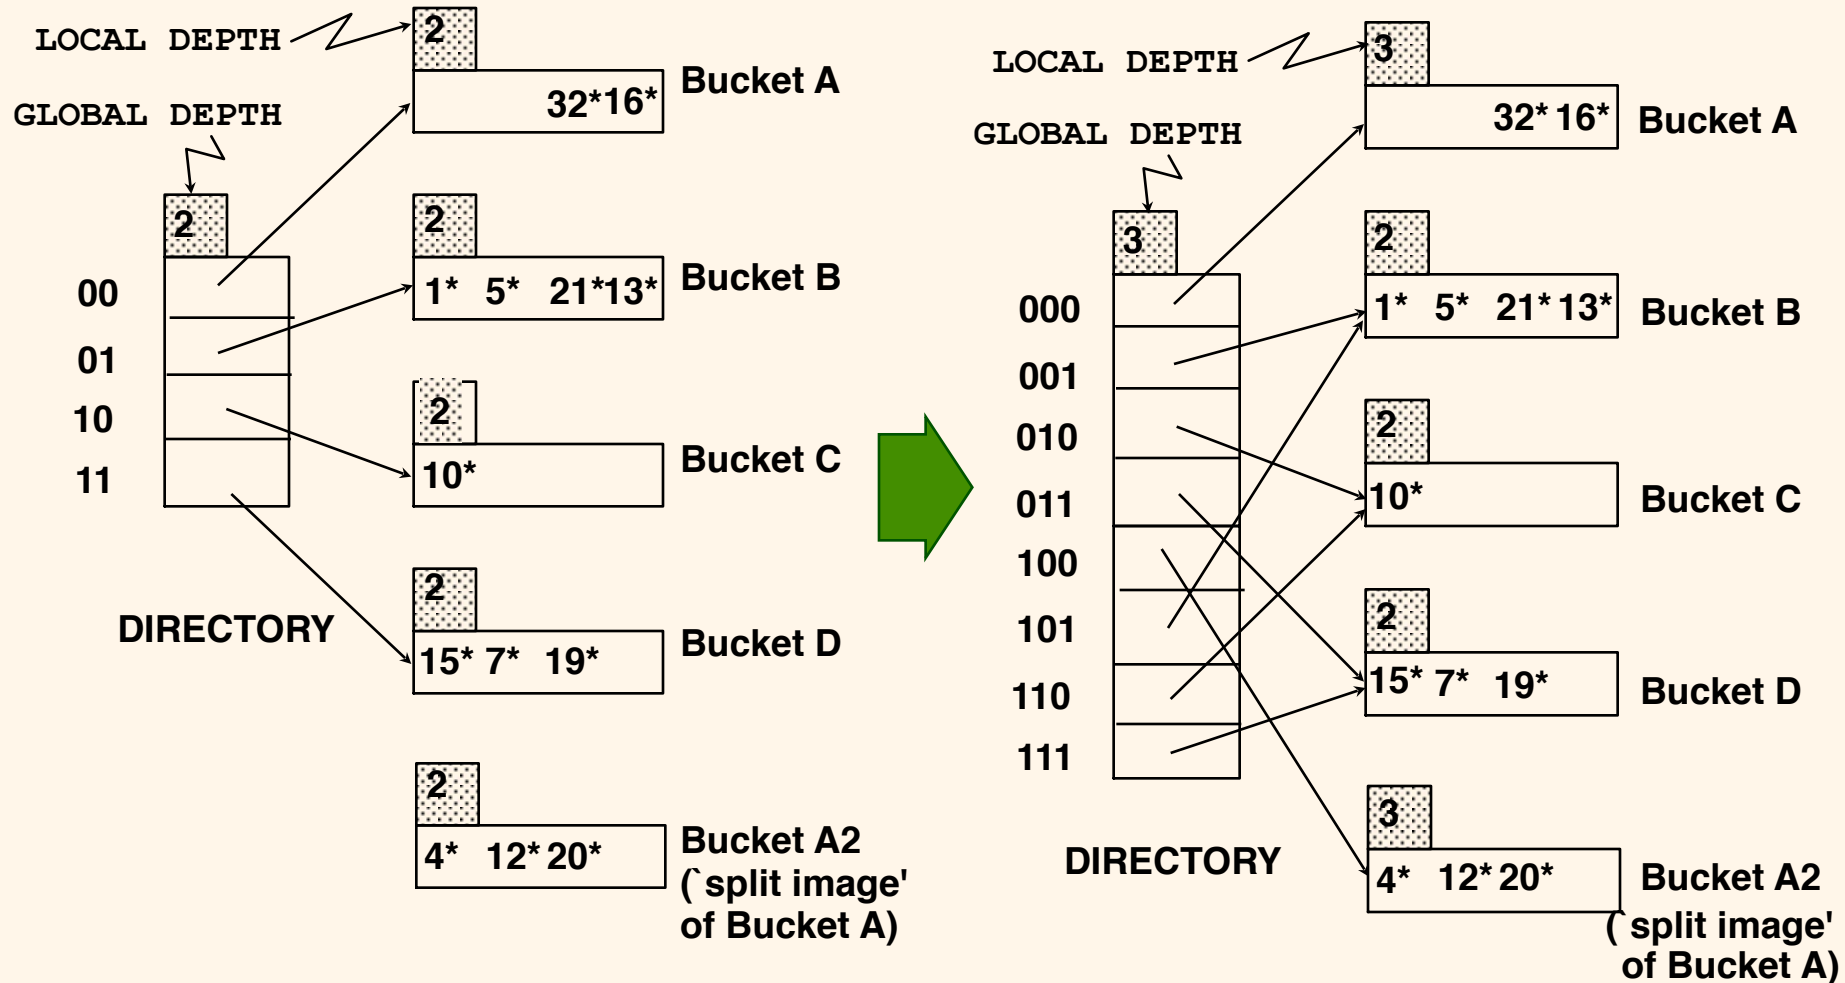
❖ Let's split that bucket (and only that one) into 2

❖ How to distribute entries among new buckets? Look at 3rd bit from the right!





Insert $h(r)=20$ (Causes Doubling)





Global and local depth

- ❑ *Global depth of directory:* Max # of bits needed to tell which bucket an entry belongs to.
- ❑ *Local depth of a bucket:* # of bits used to determine if an entry belongs to this bucket.
 - Bucket C has local depth 2, so all entries will share the same last 2 bits but may have a different 3rd-to-last one!



Points to Note

- ❑ Splitting does not always require directory doubling!
 - E.g. if we now insert lots of values into Bucket C and need to split it, no need to double directory
- ❑ When does bucket split cause directory doubling?
 - When its local depth pre-split was equal to the global depth!



What about deletions?

- ❑ If removal of data entry makes bucket empty, can be merged with 'split image'.
- ❑ If each directory element points to same bucket as its split image, can halve directory.
- ❑ In practice may omit this step and leave index as is
 - Relatively common for other indexes too e.g. B+ trees
 - Justification: deletes typically less frequent than inserts




Performance of Extendible Hashing

- ❑ If directory fits in memory, equality search answered with one disk access; else two.
- ❑ Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large.



Performance of Extendible Hashing

- ❑ Doubling the directory is a "stop the world" operation and may still take a while
- ❑ May still need overflow pages if we have **hash collisions** (two data entries/search keys have *same hash value*)




Linear Hashing

- ❑ Another dynamic hashing scheme, an alternative to Extendible Hashing.
- ❑ LH is similar to EH but does not use a directory
 - Allows some (more) overflow pages instead



Linear Hashing

- ❑ Directory avoided in LH by using overflow pages, and choosing bucket to split **round-robin**.
- ❑ Do not necessarily split the bucket that has overflowed, but rather the bucket whose "turn" it is.




Linear Hashing

- ❑ Splitting proceeds in rounds. Round ends when all N_R initial (for round R) buckets are split.
- ❑ Current round number is *Level*.



Overview of LH File

❑ In the middle of a round.





The big question

- ❑ How do we quickly find the right bucket for search key k ?
- ❑ After all, that's the whole reason we are doing hashing...
- ❑ Linear hashing: have TWO hash functions in play at all times




Hash function families

- ☐ Use a family of hash functions: h_0, h_1, h_2
- ☐ Range of h_{i+1} is twice the range of h_i
- ☐ Can generate them from some "starter" function h that maps search keys to integers
- ☐ Say initial number of buckets is $N = 2^d$
- ☐ Can take h_0 as last d bits of $h(k)$
- ☐ And h_i is last $d+i$ bits of $h(k)$
- ☐ E.g. suppose $N = 4...$




Why we don't need a directory





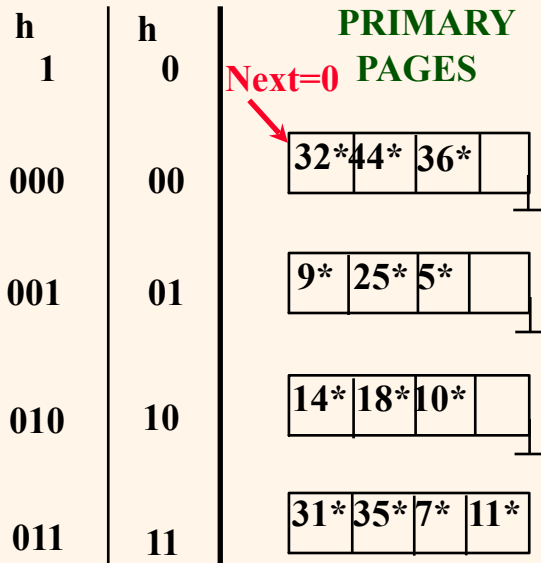
Why we don't need a directory

- ☐ For search, need to know which *Level* (round) we are in.
- ☐ To find bucket for data entry r , find $\mathbf{h}_{Level}(r)$:
 - If $Next \leq \mathbf{h}_{Level}(r) \leq N_{R'}$, r belongs here.
 - Else, r could belong to bucket $\mathbf{h}_{Level}(r)$ or bucket $\mathbf{h}_{Level}(r) + N_{R'}$; must apply $\mathbf{h}_{Level+1}(r)$ to find out.




Example – inserting 43*

Level=0, N=4

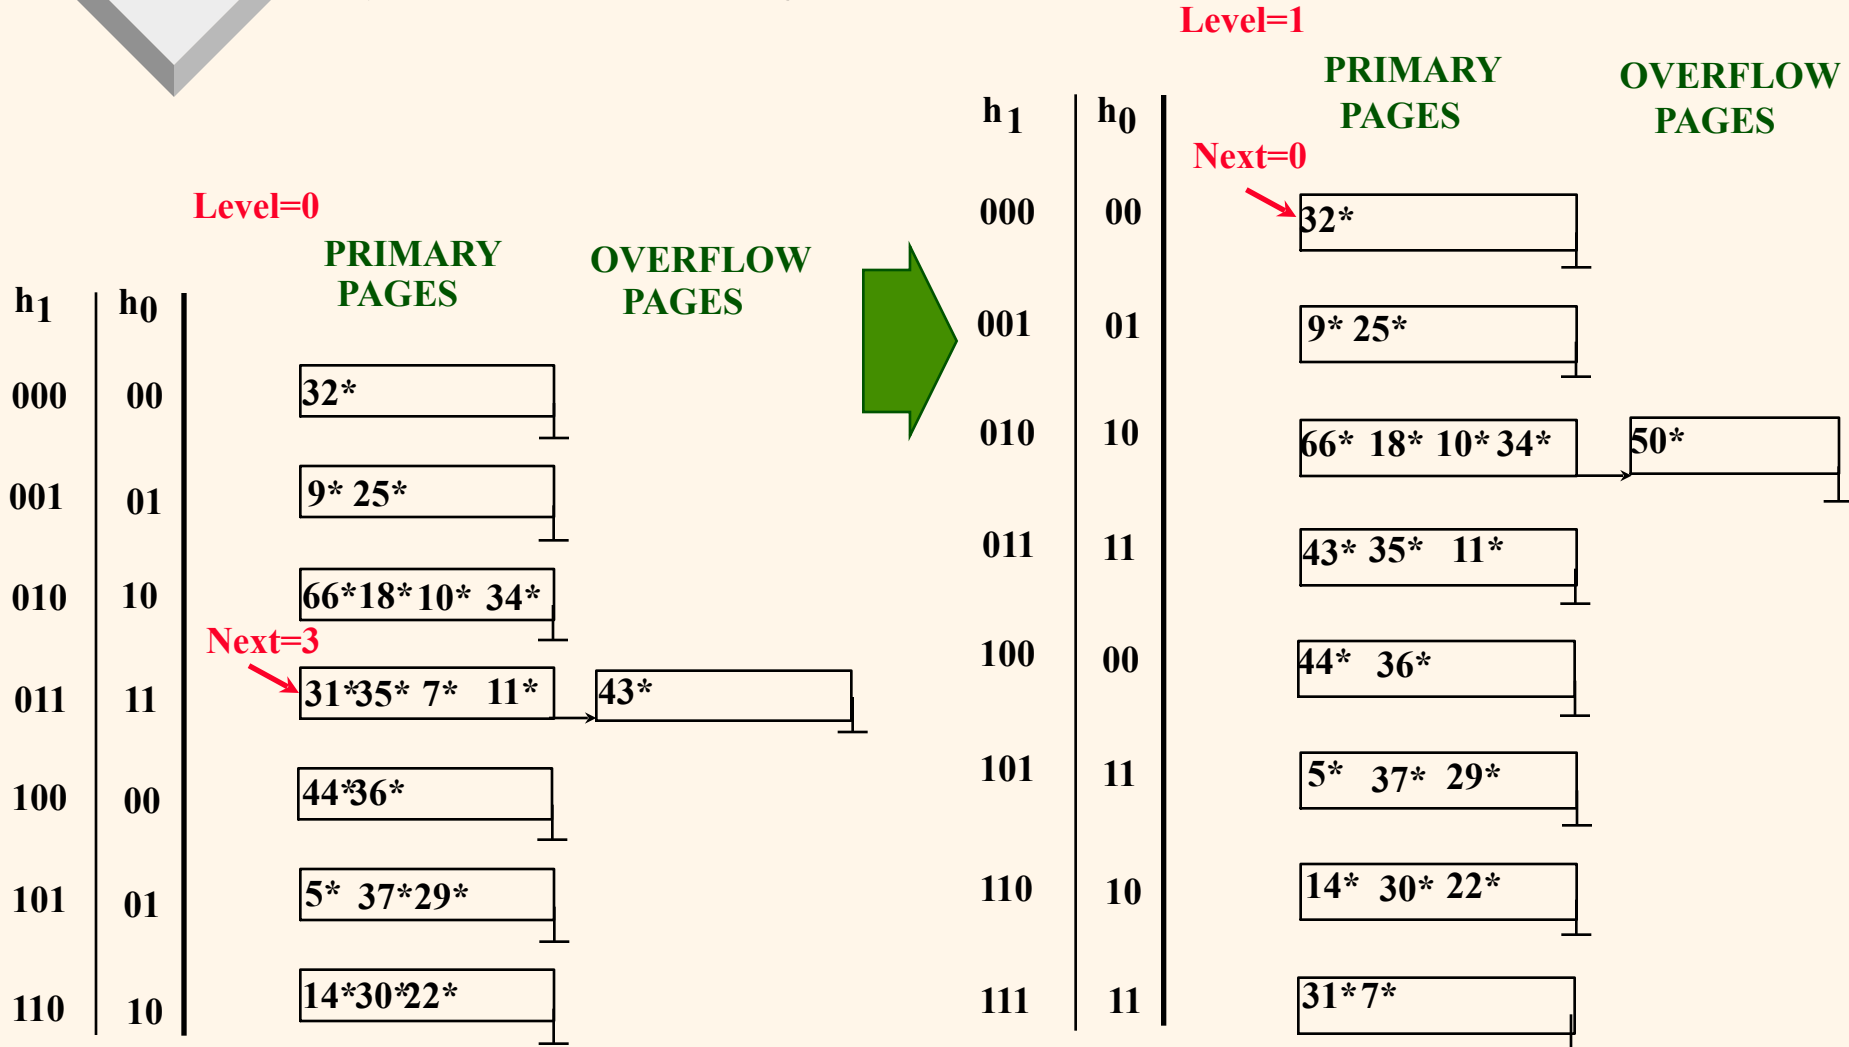


Level=0



Note this is NOT a directory

Example: End of a Round, insert 50*





Linear Hashing Continued

- ❑ Insert: Find bucket by applying \mathbf{h}_{Level} or $\mathbf{h}_{Level+1}$:
 - If bucket to insert into is full:
 - ❑ Add overflow page and insert data entry.
 - ❑ (Maybe) Split *Next* bucket and increment *Next*.
- ❑ Can choose any criterion to trigger split, eg. desired occupancy
- ❑ Similarities/ differences to doubling directory in EH



Linear Hashing Continued

- ❑ Delete: can be implemented as reverse of insert, or just ignored (leave index as is)
- ❑ Cost of lookup: 1 I/O if primary bucket pages consecutive on disk



Summary

- ❑ Hash-based indexes: best for equality searches, cannot support range searches.
- ❑ Static Hashing can lead to long overflow chains.



Summary

- ❑ Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it.
 - Directory to keep track of buckets, doubles periodically.
 - Can get large with skewed data; additional I/O if this does not fit in main memory.



Summary (Contd.)

- ☐ Linear Hashing avoids directory by splitting buckets round-robin, and using overflow pages.
- Space utilization could be lower than Extendible Hashing, since splits not concentrated on 'dense' data areas.
 - Can tune criterion for triggering splits to trade-off slightly longer chains for better space utilization.